

DCE — what to gather, and why each one is needed

August 17, 2026

CONTENTS

DCE — what to gather, and why each one is needed	
The rule, before anything else	
1. Real KYC bundles — the highest-value item	
2. Mixed documents — part text, part scan	
3. Multi-page scans, for the Azure DI latency question	
4. Two documents on ONE page — a gap, and a decision for you	
5. Photo ID specimens — the original list, unchanged	
Priority	
Why these three matter more than the rest	
What makes a specimen usable	
What I do NOT need	
Where to put them	

DCE — what to gather, and why each one is needed

For: Nilendu · 17 Aug 2026 Supersedes the earlier photo-ID list. Scope is **US, Canada, Mexico** — the Indian doctypes were removed and live on the `archive/india-doctypes` branch.

Everything below is ordered by **what it unblocks**, not by how easy it is to find. Items 1 and 2 unblock measurements that cannot be made at all today; item 5 is the only one that was on the original list.

The rule, before anything else

Only specimens carrying fabricated data. Never a real person's document — not yours, not a colleague's, not one found by image search.

Every identity document here is regulated personal data (US state ID laws, PIPEDA in Canada, LFPDPPP in Mexico). A real card in a test corpus is a breach sitting in a git history, and this has already happened once: three files were deleted from this repo for carrying real individuals' identifiers.

Acceptable sources, best first:

1. **Issuer specimen** — the “what the new card looks like” page. USCIS, IRCC, INE and several US DMVs publish these.
2. **Synthetic sample** — a mock-up made for training or fraud-awareness material.
3. **Redacted-and-fabricated** — a real layout with every value replaced. Keep the *printed issuer wording* intact; that is what the classifier reads. The values only matter for testing extraction.

Drop files in `corpus/<cc>/` as `<doctype_id>.pdf|png|jpg`, one line per file in `corpus/<cc>/manifest.jsonl`.

1. Real KYC bundles — the highest-value item

Why: segmentation ships today measured only against *synthetic* joins of corpus files. Boundary precision is 100% (every split it proposes is real), but recall is **35%**, and on same-size all-text documents with no form number it is **0%**. The corpus contains no real bundle at all, so the number that decides whether segmentation is useful to you has never been measured on the shape you actually receive.

What to gather — 3 to 5 files:

Shape	What it proves
a Typed cover letter/form + scanned attachments behind it	the mixed-read path end to end
b 2–3 photographed IDs concatenated into one PDF	boundary recall on the commonest KYC shape
c A bundle where two documents are the same page size, both typed, neither carrying a form number	the known 0%-recall blind spot, on real material
d An emailed bundle (<code>.eml</code> / <code>.msg</code>) with attachments	a path with no coverage at all today

For each, note in the filename or a covering note **which pages are which document** — that is the ground truth a recall measurement scores against. Page ranges are enough: `passport p1–2, utility bill p3, statement p4–6`.

2. Mixed documents — part text, part scan

Why: the silent-loss bug fixed this week (a typed page followed by photographed pages dropped the photographed ones with `status=ok` and no warning) is proven **only by unit fixtures I constructed**. The corpus contains no genuinely mixed document, so the fix has never been confirmed against a real one.

What to gather — 2 or 3:

- A form where the applicant's typed pages are followed by a photographed ID

- A statement with a scanned signature page at the end
- Anything where a page is a **photograph of a page** rather than a rendering

Also useful, and unusual: a scanned PDF that carries a *bad* OCR text layer from a previous tool. That is the case the glyph-sanity check exists for and there is no example of it anywhere in the corpus.

3. Multi-page scans, for the Azure DI latency question

Why: your first real DI call took **9.1 seconds and 16 polls for a single page**. The timeout is 30 s and the poll cap is 60, so roughly 3× headroom — **for one page**. Nobody knows yet whether that scales linearly.

What to gather: one scanned document of **5 pages** and one of **10+**. If a 10-page scan approaches 30 s, the fix is one environment variable, but it has to be discovered before a customer hits it rather than after.

Run them and send me the `ocr.done` line — `polls=` and `ms=` are the two numbers that matter.

4. Two documents on ONE page — a gap, and a decision for you

This does not work today, and I would rather tell you than have you find it.

Segmentation is **page-level**. A `Boundary` is a page number; a `Segment` is a page range. Two documents photocopied side by side onto one sheet cannot be separated, because the smallest unit the code has is a page.

Measured just now on a page carrying a W-9 and a bank statement side by side:

```
pages in file      : 1
boundaries proposed : []
segments returned  : 1
  [1-1] us_w9 conf=0.766 abstained=False
```

It returns **one confident answer** and never mentions the second document. Confident, not abstaining, which is the worse failure of the two.

What I need from you: is this shape real in your intake? Two IDs on one photocopied sheet, front and back of a card on one page, a passport bio page beside a visa page.

- **If it is rare** — leave it. Detecting it needs sub-page region analysis, which is a different and much larger piece of work than page segmentation.
 - **If it is common** — send me 3–4 examples and I will scope it properly. It is buildable (Azure DI returns paragraph bounding boxes, so the geometry is available), but it is a new capability rather than a tuning change, and it should be measured before it is promised.
-

5. Photo ID specimens — the original list, unchanged

Still needed, still in this order. These are the documents the registry knows about but the corpus has no example of.

PRIORITY

#	Document	Registry id	Where specimens exist
1	Mexican INE / credencial para votar	<code>mx_ine</code>	INE publishes specimen imagery of each model (2013 / 2019 / current)
2	US driver's licences — 3 or 4 states	<code>us_drivers_license</code>	Many DMVs publish a “your new licence” page. Prefer states with different layouts: CA, TX, NY, FL
3	US state ID (non-driver)	<code>us_state_id</code>	Same sources; the wording differs from a licence and that difference is what the classifier reads
4	Canadian PR card	<code>ca_pr_card</code>	IRCC specimen imagery
5	Canadian provincial photo ID	<code>ca_provincial_photo_id</code>	Provincial issuers
6	US passport card	<code>us_passport_card</code>	State Dept. specimen
7	Mexican matrícula consular	<code>mx_matricula_consular</code>	SRE specimen

WHY THESE THREE MATTER MORE THAN THE REST

`us_drivers_license`, `ca_drivers_license` and `ca_nexus` have **every decisive anchor gated to the title zone**. Zones only exist on the `azure_layout` reading — a text layer, local OCR and `azure_read` all produce zone-free output. So on any other path those three **cannot be decisively identified at all**.

If you gather nothing else from this section, gather driver's licences, and capture them as **images** so they go through Azure DI.

WHAT MAKES A SPECIMEN USABLE

The classifier reads **printed issuer text**, not the photograph. Values can be nonsense; the card's own wording must be legible.

- The issuing authority's name as printed
- Any form, control or document number — these are the strongest single signal there is
- The MRZ if the document has one, legible
- Field captions in the issuer's own words and language

A cropped photo of just the portrait is worthless. A slightly blurry full-card image with readable captions is valuable.

What I do NOT need

- Anything with a real person's data on it, in any form, for any reason
 - High-resolution scans — legible beats large; 200–300 dpi is plenty
 - More born-digital forms. The corpus already has 129 doctypes' worth and they are the easy case; every gap above is a **scan** or a **bundle**
-

Where to put them

Drop them anywhere convenient and tell me the path — I will place them in `corpus/`, write the manifest lines and re-run the measurements. If they are sensitive enough that they should not be committed, say so and I will keep them out of git and measure locally instead.

One note on item 1: bundles are the only item where I need you to tell me **which pages are which document**. Everything else is self-describing.